


# Deep Pepper: An Expert Iteration Based Chess Agent in the RL Setting

Sai Krishna, Kyle Goyette, Ahmad Shamseddin, Breandan Considine, Giancarlo Kerg

# Introduction/Motivation

- Computer chess widely studied for more than 40 years.
- Alpha Zero plays chess learning entirely by self-play
  - Tabula rasa shows the power of generalized AI
  - But could embedding knowledge lead to a better chess player overall?
- How can we accelerate training?
  - **Custom feature representation**
    - Substantially reduced feature representation size
    - Uses hand-crafted features developed by human experts
  - **Pretraining the networks**
    - Thanks to Stockfish
  - **Early termination of games**
    - Using Oracle (Stockfish)



# Monte Carlo Tree Search



Each edge stores:

- visit count  $\mathbf{N(s,a)}$
- Total action-value  $\mathbf{W(s,a)}$
- Mean action value  $\mathbf{Q(s,a)}$
- Prior Probabilities  $\mathbf{P(s,a)}$


## Action selection:

$$a_t = \operatorname{argmax}_a (Q(s_t, a) + U(s_t, a))$$

$$U(s, a) = c_{puct} P(s, a) \frac{\sqrt{\sum_b N(s, b)}}{1 + N(s, a)}$$

## Expand and Evaluate

- Evaluation: from value network
- Expansion: Initialize edge stats from leaf node



## Back up

$$N(s_t, a_t) = N(s_t, a_t) + 1$$
$$W(s_t, a_t) = W(s_t, a_t) + v$$
$$Q(s_t, a_t) = \frac{W(s_t, a_t)}{N(s_t, a_t)}$$

## Algorithm:

Repeat

1. Game Generation via MCTS
  - a. Save MCTS policies and game outcome for each state
2. Train network, via backpropagation
3. Test new network against last iteration

$$l = (z - v)^2 - \boldsymbol{\pi}^\top \log \mathbf{p} + c \|\boldsymbol{\theta}\|^2$$

Diagram illustrating the components of the loss function  $l$ :

- $z$ : Result at the end of each game (From value network)
- $v$ : From value network
- $\boldsymbol{\pi}^\top \log \mathbf{p}$ : Strengthened policy obtained from MCTS (From policy network)
- $c \|\boldsymbol{\theta}\|^2$ : Regularization

## This is roughly Classification-Based Approximate Policy Iteration

**Policy Improvement:** Each MCTS call strengthens the current policy (predictions in the network) and strengthened policy from MCTS is projected back into network functional space via training.

**Policy Evaluation:** True game scores using the strengthened policy. Projected into function space via network training.

- **Relation to CAPI**
- **Generalization to multiple actions**

$$\hat{L}_n^{\pi_k}(\pi) \triangleq \int_{\mathcal{X}} \mathbf{g}_{\hat{Q}^{\pi_k}}(x) \mathbb{I}\{\pi(x) \neq \operatorname{argmax}_{a \in \mathcal{A}} \hat{Q}^{\pi_k}(x, a)\} d\nu_n$$

$$\mathbf{g}_Q(x) \triangleq |Q(x, 1) - Q(x, 2)| \quad \text{for all } x \in \mathcal{X}.$$

$$\mathbb{P}_\nu (0 < \mathbf{g}_{Q^\pi}(X) \leq \varepsilon) \triangleq \int_{\mathcal{X}} \mathbb{I}\{0 < \mathbf{g}_{Q^\pi}(x) \leq \varepsilon\} d\nu(x) \leq c_g \varepsilon^\zeta.$$

**Algorithm** CAPI( $\Pi, \nu, K$ )

**Input:** Policy space  $\Pi$ , State distribution  $\nu$ , Number of iterations  $K$

**Initialize:** Let  $\pi_{(0)} \in \Pi$  be an arbitrary policy

**for**  $k = 0, 1, \dots, K - 1$  **do**

Construct a dataset  $\mathcal{D}_n^{(k)} = \{X_i\}_{i=1}^n$ ,  $X_i \stackrel{\text{i.i.d.}}{\sim} \nu$

$\hat{Q}^{\pi_k} \leftarrow \text{PolicyEval}(\pi_k)$

$\pi_{k+1} \leftarrow \operatorname{argmin}_{\pi \in \Pi} \hat{L}_n^{\pi_k}(\pi)$  (action-gap-weighted classification)

**end for**

- **Experimental Results:**
  - **Trained network for 100 games:**
    - **Wins 76% of the time against a pretrained model**
  - **Pretrained model**
    - **Wins 64% of games against randomly initialized model**
- **Future Work**
  - **Approximate ELO rating**
  - **More training**
  - **Using opening books**
    - **Custom opening book from online databases**
  - **Gradual shift from MCTS to alpha-beta pruning**
  - **Parallelize MCTS algorithm to speed up gameplay**

Merci Beaucoup